

# Elements of Data Science - F20

---

## Final Review

---

This is intended as a guide and is not guaranteed to be inclusive.

Material considered fair for the exam is anything from class and slides.

## Data Science Tools

- Data Science workflow
- Jupyter+Ipython Notebooks
- conda Virtual Environments
- using Git to pull code and materials

## Python Intro/Review Numpy and Pandas

- Importing modules
- Defining functions
- String Formatting
- What are Exceptions?
- Using assert
- Basic Python data types
- Collections module: Counter, defaultdict
- Python flow control: if: elif: else: , for x in xs:
- Sorting with lambda functions as the key
- List Comprehensions
- Numpy
  - arrays
  - indexing/slicing
  - Boolean masks and bitwise operations
- Pandas
  - Series
  - DataFrames
  - indexing/slicing
  - .describe

## Visualization and Data Exploration

- Matplotlib
  - plotting using matplotlib
  - using plt.subplots()
  - modifying plots using ax
- Variable Types
- Central tendencies
  - mean
  - median
- Spread
  - variance
  - std deviation
  - skew

- IQR
- Correlation
  - Pearson Correlation Coefficient
- Univariate Plotting
  - hist
  - seaborn.histplot()
  - boxplots
- Bivariate Plotting
  - scatter
  - jointplot
  - pairplot

## Hypothesis Testing

- Random Sampling vs Population Distribution
- Sample Statistic
- Confidence Intervals
- Normal (Gaussian) Distribution
  - Standard Normal Distribution
  - Z-Score
- Central Limit Theorem
- Bootstrap Sampling
- A/B Test
- Hypothesis Testing
  - Type I and II error
  - Significance and Power
  - Permutation Tests
  - One-tailed vs Two-tailed
  - p-values
- Calculating “How many observations?”
  - what 4 values are related?
- Multi-Armed Bandit
  - benefits of using
  - greedy
  - epsilon-greedy

## Intro to ML

- Dimensions of ML
  - Interpretation vs Prediction
  - Learning Paradigms (SL,UL,etc.)
  - Regression vs Classification
  - Binary, Multiclass, Multilabel
- sklearn common functions
  - .fit
  - .predict
  - .predict\_proba

## Machine Learning Models

- Simple Linear Regression
  - Residuals in linear models

- Interpreting Coefficients of OLS
  - Colinearity
  - Multiple Linear Regression
  - Logistic Regression
  - Concept of Gradient Descent
  - SVM
  - One vs. Rest for Multiclass/Multilabel
  - k-Nearest Neighbor
  - Decision Trees
  - Ensembles
    - Random Forest
    - Gradient Boost
    - Stacking
- 

## After the Midterm

---

### Model Evaluation

- Generalization
  - Train/Test split
  - stratification
- Overfitting/Underfitting
  - Bias/Variance Tradeoff
- Baseline Models
- Tuning Hyperparameters and Model Selection
  - k-Fold Cross Validation
  - Grid Search
- Plotting Model Fit
  - Validation Curve
- Metrics: Classification
  - Confusion Matrix
  - Accuracy/Error
  - Precision
  - Recall
  - Precision-Recall Curve
  - F1 Score
  - ROC Curve
  - ROC AUC
- Metrics: Regression
  - $R^2$
  - Mean Squared Error
  - RMSE
- Regularization
  - Ridge
  - LASSO
  - ElasticNet

### Data Cleaning

- Duplicates
- Missing Data

- Dummy Variables
- Rescaling
- Dealing With Skew
- Removing Outliers

## Feature Engineering

- Binning
- One-Hot encoding
- Derived Features

## Dimensionality Reduction

- Feature Selection
  - LASSO
  - Tree Based Models Feature Importance
  - Univariate Tests
  - Recursive Feature Selection
- Feature Extraction
  - PCA

## NLP and Topic Modeling

- What is a corpus?
- Tokens and Tokenization
- Vocabulary
- Bag Of Words representation
- n-grams
- Term Frequency
- Document Frequency
- Stopwords
- TfIdf
- Latent Dirichlet Allocation (general concept)
  - per document topic distribution
  - per topic term distribution

## Clustering

- k-Means
- Hierarchical Agglomerative Clustering
  - linkage

## Recommendation Engines

- Content-Based Filtering
- User-Based Collaborative Filtering
- Issues
- Evaluating

## Timeseries

- unique characteristics of timeseries data
- timeseries in pandas
- indexing with a DatetimeIndex

- converting data to datetime with `pd.to_datetime()`
- Shifting
- Resampling/Frequencies
- Upsampling vs Downsampling
- Moving/Rolling Window functions

## **Dealing with Imbalanced Data**

- Oversampling minority class
- Undersampling majority class
- SMOTE and ADASYN (general concept)

## **Data Processing (ETL and API)**

- What does ETL stand for?
- Difference between csv and json
- sklearn Pipelines
  - ColumnTransformer
- What can we use the flask python library for?