

Elements Of Data Science - F2020

Week 8: Data Cleaning and Feature Engineering

11/9/2020

TODOs

- Readings:
 - PDSH 5.9 PCA
 - [Recommeded] PML Ch 8
- HW2, Due Thurs Nov. 12th, 11:59pm
- Answer and submit Quiz 8, **Sunday Nov 15th, 11:59pm ET**

Today

- **Data Cleaning**
 - Duplicates
 - Missing Data
 - Dummy Variables
 - Rescaling
 - Dealing With Skew
 - Removing Outliers
- **Feature Engineering**
 - Binning
 - One-Hot encoding
 - Derived
 - Unstructured Data: Natural Language Processing

Questions?

Data Cleaning

Why do we need clean data?

- Want one row per observation (remove duplicates)
- Most models cannot handle missing data (remove/fill missing)
- Most models require fixed length feature vectors (engineer features)
- Different models require different types of data (transformation)
 - Linear models: real valued features with similar scale
 - Distance based: real valued features with similar scale
 - Tree based: can handle real and categorical

Example Data and `.info()`

```
In [2]: # read in example data
df_shop = pd.read_csv('../data/flowershop_data_with_dups.csv',
                      header=0,
                      parse_dates=['purchase_date'],
                      delimiter=',')

df_shop.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1001 entries, 0 to 1000
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   purchase_id     1001 non-null  int64
1   lastname        1001 non-null  object
2   purchase_date   1001 non-null  datetime64[ns]
3   stars           1001 non-null  int64
4   price           979 non-null   float64
5   favorite_flower  823 non-null   object
dtypes: datetime64[ns](1), float64(1), int64(2), object(2)
memory usage: 47.0+ KB
```

Duplicated Data

- Only drop duplicates if you know data should be unique
 - Example: if there is a unique id per row

```
In [3]: df_shop.duplicated().iloc[:3] # are first 3 rows duplicates?
```

```
Out[3]: 0    False
        1    False
        2    False
        dtype: bool
```

```
In [4]: df_shop[df_shop.duplicated(keep='first')] # default keep first duplicated row
```

```
Out[4]:
```

	purchase_id	lastname	purchase_date	stars	price	favorite_flower
1000	1010	FERGUSON	2017-05-04	2	21.0183	daffodil

```
In [5]: df_shop[df_shop.duplicated(keep=False)] # keep=False to show all duplicated rows
```

```
Out[5]:
```

	purchase_id	lastname	purchase_date	stars	price	favorite_flower
10	1010	FERGUSON	2017-05-04	2	21.0183	daffodil
1000	1010	FERGUSON	2017-05-04	2	21.0183	daffodil

Duplicated Data for Subset of Columns

In [6]: `df_shop[df_shop.duplicated(subset='purchase_id', keep=False)].sort_values(by='purchase_id')`

Out[6]:

	purchase_id	lastname	purchase_date	stars	price	favorite_flower
10	1010	FERGUSON	2017-05-04	2	21.018300	daffodil
1000	1010	FERGUSON	2017-05-04	2	21.018300	daffodil
100	1101	WEBB	2017-07-13	2	8.004356	iris
101	1101	BURKE	2017-08-16	4	18.560260	daffodil

Dropping Duplicates

```
In [7]: df_new = df_shop.drop_duplicates(subset=None      # consider subset of columns
                                     , keep='first'      # or 'last' or False)
                                     , inplace=False)
```

```
In [8]: # or can use inplace to change the original dataframe
df_shop.drop_duplicates(subset=None, keep='first', inplace=True)
```

```
In [9]: # drop rows with duplicates considering only a subset of columns
df_shop = df_shop.drop_duplicates(subset=['purchase_id'])
```

Missing Data

- Reasons for missing data
 - Sensor error (random?)
 - Data entry error (random?)
 - Survey-subject decisions (non-random?)
 - etc.
- Dealing with missing data
 - Drop rows
 - Impute from data in the same column
 - Infer from other features
 - Fill with adjacent data

Missing Data in Pandas: `np.nan`

- Missing values represented by `np.nan` : Not A Number

```
In [10]: # Earlier, we saw missing values in the dataframe summary
df_shop.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 999 entries, 0 to 999
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   purchase_id     999 non-null   int64
1   lastname        999 non-null   object
2   purchase_date   999 non-null   datetime64[ns]
3   stars           999 non-null   int64
4   price           977 non-null   float64
5   favorite_flower  821 non-null   object
dtypes: datetime64[ns](1), float64(1), int64(2), object(2)
memory usage: 54.6+ KB
```

```
In [11]: np.nan == np.nan
```

```
Out[11]: False
```

.isna() and .notna()

```
In [12]: df_shop.loc[20:21, 'price']
```

```
Out[12]: 20      NaN
         21    10.525912
         Name: price, dtype: float64
```

```
In [13]: # isna returns True where data is missing, False otherwise
         df_shop.loc[20:21, 'price'].isna()
```

```
Out[13]: 20     True
         21    False
         Name: price, dtype: bool
```

```
In [14]: # notna returns True where data is NOT missing, False otherwise
         df_shop.loc[20:21, 'price'].notna()
```

```
Out[14]: 20     False
         21     True
         Name: price, dtype: bool
```

```
In [15]: # find rows where price is missing
         df_shop[df_shop.price.isna()].head()
```

```
Out[15]:
```

	purchase_id	lastname	purchase_date	stars	price	favorite_flower
20	1020	CLARK	2017-01-05	3	NaN	NaN
41	1041	PETERS	2017-02-01	4	NaN	orchid
54	1054	GREEN	2017-02-13	5	NaN	daffodil
63	1063	BARNETT	2017-08-27	4	NaN	gardenia
145	1145	CARROLL	2017-07-29	3	NaN	tulip

Missing Data: Drop Rows

```
In [16]: df_shop.shape
```

```
Out[16]: (999, 6)
```

```
In [17]: # drop rows with nan in any column  
df_shop.dropna().shape
```

```
Out[17]: (801, 6)
```

```
In [18]: # drop only rows with nan in price using subset  
df_shop.dropna(subset=['price']).shape
```

```
Out[18]: (977, 6)
```

```
In [19]: # drop only rows with nans in all columns  
df_shop.dropna(how='all').shape
```

```
Out[19]: (999, 6)
```

Missing Data: Drop Rows Cont.

```
In [20]: # save a new dataframe with dropped rows
df_new = df_shop.dropna()
df_new.shape
```

Out[20]: (801, 6)

```
In [21]: # drop rows in current dataframe
df_new = df_shop.copy()
df_new.dropna(inplace=True)
df_new.shape
```

Out[21]: (801, 6)

```
In [22]: # How many total nan's?
df_shop.isna().sum().sum()
```

Out[22]: 200

Missing Data: Drop Rows

- Pros:
 - easy to do
 - simple to understand
- Cons:
 - potentially large data loss

Missing Data: Impute

- Impute: fill in with some data
- Use `.fillna()`
- Common filler values:
 - mean
 - median
 - "most frequent"
 - 0

Missing Data: Impute

```
In [23]: df_shop.price.mean()
```

```
Out[23]: 23.408197893394266
```

```
In [24]: # make a copy to keep our original df
df_new = df_shop.copy()
```

```
In [25]: # fill missing price with mean of price
df_new.price = df_shop.price.fillna(df_shop.price.mean())
```

```
In [26]: # check to make sure all nulls filled
sum(df_new.price.isna())
```

```
Out[26]: 0
```

```
In [27]: # inplace works here as well
df_new.price.fillna(df_new.price.mean(),inplace=True)
```

```
In [28]: # can also handle categorical data
df_new.favorite_flower.fillna(df_shop.favorite_flower.mode().iloc[0],inplace=True)
df_new.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 999 entries, 0 to 999
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   purchase_id     999 non-null   int64
1   lastname        999 non-null   object
2   purchase_date   999 non-null   datetime64[ns]
3   stars           999 non-null   int64
```

Missing Data: Impute

- Pros:
 - easy to do
 - simple to understand
- Cons:
 - may miss feature interactions

Missing Data: Infer

- Predict values of missing features using a model
- Ex: Can we predict price based on any of the other features?
- Additional feature engineering may be needed prior to this

```
In [29]: from sklearn.linear_model import LinearRegression

df_tmp = df_shop.copy()

not_missing = df_tmp.price.notna()
missing = df_tmp.price.isna()

lr = LinearRegression().fit(df_tmp.loc[not_missing, ['stars']],
                           df_tmp[not_missing].price)

df_tmp.loc[missing, 'price'] = lr.predict(df_tmp.loc[missing, ['stars']])
```

Missing Data: Adjacent Data

- Use `.fillna()` with method:
 - `ffill`: propagate last valid observation forward to next valid
 - `bfill`: use next valid observation to fill gap backwards
- Use when there is reason to believe data not i.i.d. (eg: timeseries)

```
In [30]: df_shop.price.loc[19:21]
```

```
Out[30]: 19    20.451789
         20         NaN
         21    10.525912
         Name: price, dtype: float64
```

```
In [31]: df_shop.price.fillna(method='ffill').loc[19:21]
```

```
Out[31]: 19    20.451789
         20    20.451789
         21    10.525912
         Name: price, dtype: float64
```

Missing Data: Dummy Columns

- Data may be missing for a reason!
- Capture "missing" before filling

```
In [32]: # storing a column of 1:missing, 0:not-missing
df_new['price_isnull'] = df_shop.price.isna().astype(int)

# can now fill missing values
df_new['price'] = df_shop.price.fillna(df_shop.price.mean())
```

Rescaling

- Often need features to be in the same scale
- Methods of rescaling
 - Standardization (z-score)
 - Min-Max rescaling
 - others...

```
In [33]: # load taxi data
taxi = pd.read_csv('../data/yellow_tripdata_2017-01_subset10000rows.csv',
                  parse_dates=['tpep_pickup_datetime', 'tpep_dropoff_datetime'])

# create trip_duration
taxi['trip_duration'] = (taxi.tpep_dropoff_datetime - taxi.tpep_pickup_datetime).dt.seconds

# select subset
df_taxi = taxi[(taxi.trip_duration < 3600) & (taxi.tip_amount > 0) & (taxi.tip_amount < 10)]
```

```
In [34]: df_taxi[['trip_duration', 'tip_amount']].agg(['mean', 'std', 'min', 'max'], axis=0)
```

```
Out[34]:
```

	trip_duration	tip_amount
mean	765.030683	2.405944
std	496.831608	1.552848
min	2.000000	0.010000
max	3556.000000	9.990000

Rescaling: Standardization

- rescale to 0 mean, standard deviation of 1
 - $X_{\text{scaled}} = (X - X.\text{mean}()) / X.\text{std}()$

```
In [35]: from sklearn.preprocessing import StandardScaler

# instantiate
ss = StandardScaler()

# fit to the data
ss.fit(df_taxi[['trip_duration', 'tip_amount']])

# transform the data
X = ss.transform(df_taxi[['trip_duration', 'tip_amount']])
X[:2]
```

```
Out[35]: array([[ -0.50127786, -0.48040987],
                [-0.16512088, -0.90546941]])
```

```
In [36]: df_new = pd.DataFrame(X, columns=['trip_duration_scaled', 'tip_amount_scaled'])
df_new.agg(['mean', 'std', 'min', 'max'], axis=0)
```

```
Out[36]:
```

	trip_duration_scaled	tip_amount_scaled
mean	4.622808e-17	-1.358307e-16
std	1.000080e+00	1.000080e+00
min	-1.535917e+00	-1.543059e+00
max	5.617987e+00	4.884357e+00

Rescaling: Min-Max

- rescale values between 0 and 1
- $X_{\text{scaled}} = (X - X_{\text{min}}()) / (X_{\text{max}}() - X_{\text{min}}())$
- removes negative values

```
In [37]: from sklearn.preprocessing import MinMaxScaler

X = MinMaxScaler().fit_transform(df_taxi[['trip_duration', 'tip_amount']])

df_new = pd.DataFrame(X, columns=['trip_duration_scaled', 'tip_amount_scaled'])
df_new.agg(['mean', 'std', 'min', 'max'])
```


Out[37]:

	trip_duration_scaled	tip_amount_scaled
mean	0.214696	0.240075
std	0.139795	0.155596
min	0.000000	0.000000
max	1.000000	1.000000

Dealing with Skew


- Many models expect "normal", symmetric data (ex: linear models)
- Highly skewed: tail has larger effect on model (outliers?)
- Transform with `log` or `sqrt`

```
In [38]: fig,ax = plt.subplots(1,3,figsize=(16,4))
sns.histplot(x=df_taxi.total_amount, ax=ax[0]);
sns.histplot(x=df_taxi.total_amount.apply(np.sqrt), ax=ax[1]); ax[1].set_xlabel('sqrt transform');
sns.histplot(x=df_taxi.total_amount.apply(np.log), ax=ax[2]); ax[2].set_xlabel('log transform');
```



Outliers

- Similar to missing data:
 - human data entry error
 - instrument measurement errors
 - data processing errors
 - natural deviations





Outliers

- Why worry about them?
 - can give misleading results
 - can indicate issues in data/measurement
- Detecting Outliers
 - understand your data!
 - visualizations
 - $1.5 \times \text{IQR}$
 - z-scores
 - etc..

Detecting Outliers

```
In [39]: df = pd.DataFrame(np.random.normal(50,20,1000), columns=['measure'])
df = df.append(pd.DataFrame(np.random.normal(120,1,20), columns=['measure']))

fig,ax = plt.subplots(1,2, figsize=(14,4))
sns.histplot(x=df.measure,ax=ax[0]);
sns.boxplot(x=df.measure,ax=ax[1]);
```




Detecting Outliers with z-score

```
In [40]: # zscore
df['measure_zscore'] = (df.measure - df.measure.mean()) / df.measure.std()

fig, ax = plt.subplots(1,3,figsize=(16,4))
sns.histplot(x=df.measure,ax=ax[0]);
sns.histplot(x=df.measure_zscore, ax=ax[1]);

keep_idx = np.abs(df.measure_zscore) < 2
sns.histplot(x=df[keep_idx].measure_zscore, ax=ax[2]);
```



Other Outlier Detection Methods

- Many more parametric and non-parametric methods
 - Standardized Residuals
 - DBScan
 - EllipticEnvelope
 - IsolationForest
 - other Anomaly Detection techniques
 - See [sklearn docs on Outlier Detection](#) for more details

Dealing with Outliers

- How to deal with outliers?
 - drop data
 - treat as missing
 - encode with dummy variable first?

Data Cleaning Review

- duplicate data
- missing data
- rescaling
- dealing with skew
- outlier detection

Feature Engineering

- Binning
- One-Hot encoding
- Derived

Binning

- Transform continuous features to categorical
- Use:
 - `pd.cut`
 - `sklearn.preprocessing.KBinsDiscretizer` (combined binning and one-hot-encoding)

```
In [41]: trip_duration_bins = [df_taxi.trip_duration.min(),
                              df_taxi.trip_duration.median(),
                              df_taxi.trip_duration.quantile(0.75),
                              df_taxi.trip_duration.max(),
                              ]
```

```
In [42]: df_new = df_taxi.copy()
df_new['trip_duration_binned'] = pd.cut(df_taxi.trip_duration,
                                       bins=trip_duration_bins,           # can pass bin edges or number of bins
                                       labels=['short', 'medium', 'long'],
                                       right=True,                         # all bins right-inclusive
                                       include_lowest=True,                 # first interval left-inclusive
                                       )
df_new[['trip_duration', 'trip_duration_binned']].iloc[:4]
```

```
Out[42]:
```

	trip_duration	trip_duration_binned
1	516	short
2	683	medium
7	834	medium

One-Hot Encoding

- Encode categorical features for models that can't handle categorical (eg. Linear)
- One column per category, '1' in only one column per row
- Use `pd.get_dummies()` or `sklearn.preprocessing.OneHotEncoder`

```
In [43]: pd.get_dummies(df_new.trip_duration_binned, prefix='trip_duration').iloc[:3]
```

```
Out[43]:
```

	trip_duration_short	trip_duration_medium	trip_duration_long
1	1	0	0
2	0	1	0
7	0	1	0

```
In [44]: df_new.join(pd.get_dummies(df_new.trip_duration_binned, prefix='trip_duration')).iloc[:3, -5:]
```

```
Out[44]:
```

	trip_duration	trip_duration_binned	trip_duration_short	trip_duration_medium	trip_duration_long
1	516	short	1	0	0
2	683	medium	0	1	0
7	834	medium	0	1	0

```
In [45]: pd.get_dummies(df_new).iloc[:3, -7:]
```

```
Out[45]:
```

	total_amount	trip_duration	store_and_fwd_flag_N	store_and_fwd_flag_Y	trip_duration_binned_short	trip_duration_binned_medium	trip_duration_binned_long
1	9.96	516	1	0	1	0	0
2	10.30	683	1	0	0	1	0
7	16.64	834	1	0	0	1	0

One-Hot Encoding with sklearn

```
In [46]: from sklearn.preprocessing import OneHotEncoder

ohe = OneHotEncoder(categories=[['short', 'medium', 'long']], # or leave as 'auto'
                    sparse=False,
                    handle_unknown='ignore') # will raise error otherwise

ohe.fit(df_new[['trip_duration_binned']])
ohe.categories_
```

```
Out[46]: [array(['short', 'medium', 'long'], dtype=object)]
```

```
In [47]: ohe.transform(df_new[['trip_duration_binned']])[:3]
```

```
Out[47]: array([[1., 0., 0.],
               [0., 1., 0.],
               [0., 1., 0.]])
```

Bin and One-Hot Encode with sklearn

```
In [48]: from sklearn.preprocessing import KBinsDiscretizer

# NOTE: We're not setting the bin edges explicitly
#       For control over bin edges, use Binarizer
kbd = KBinsDiscretizer(n_bins=3,
                      encode="onehot",      # or onehot (sparse), ordinal
                      strategy="quantile",  # or uniform or kmeans (clustering)
                      ).fit(df_new[['trip_duration']])

kbd.bin_edges_
```

```
Out[48]: array([array([2.000e+00, 4.780e+02, 8.700e+02, 3.556e+03])], dtype=object)
```

```
In [49]: df_new[['trip_duration']].head(3)
```

```
Out[49]:
```

	trip_duration
1	516
2	683
7	834

```
In [50]: kbd.transform(df_new[['trip_duration']])[:3]
```

```
Out[50]: <3x3 sparse matrix of type '<class 'numpy.float64'>'
         with 3 stored elements in Compressed Sparse Row format>
```

```
In [51]: kbd.transform(df_new[['trip_duration']])[:3].todense()
```

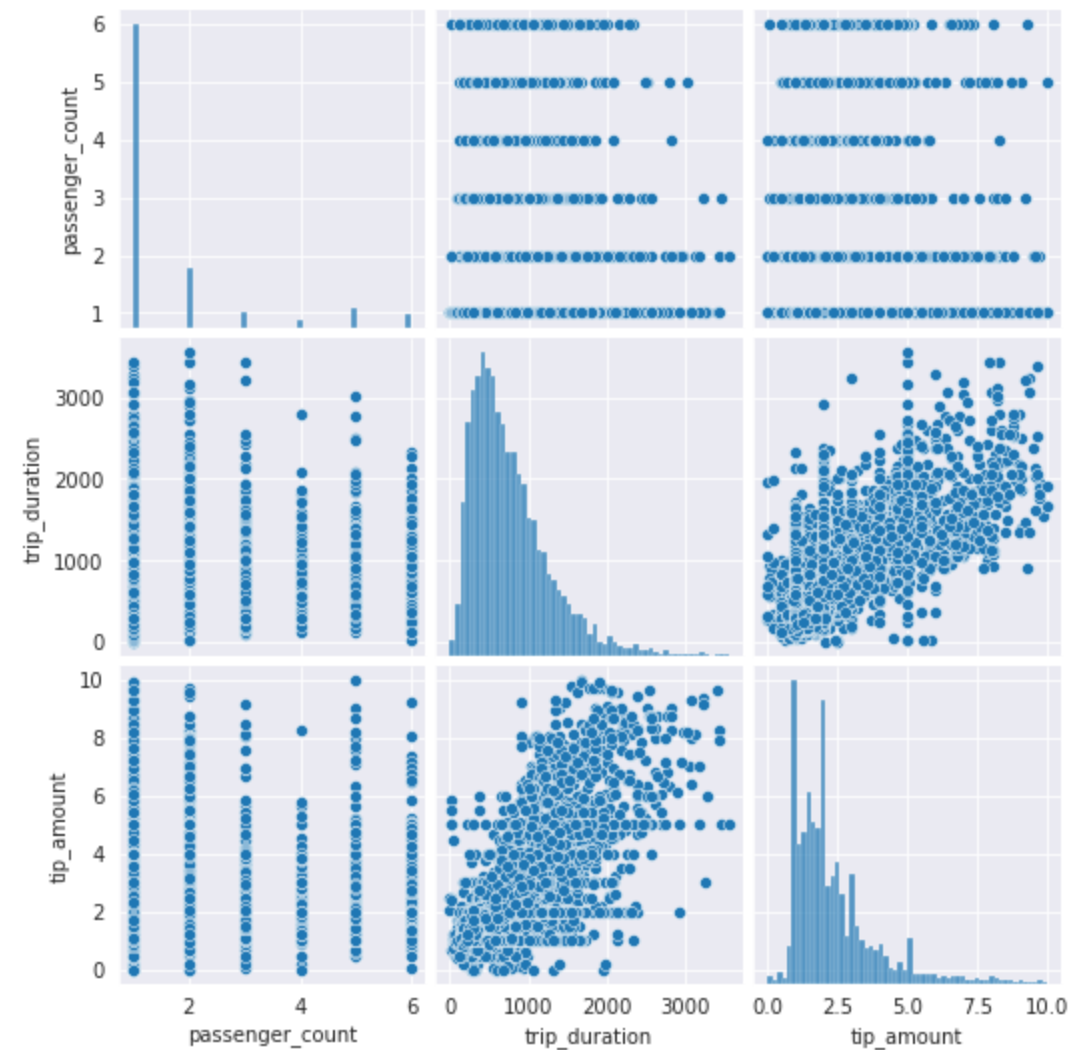
```
Out[51]: matrix([[0., 1., 0.],
                 [0., 1., 0.],
                 [0., 1., 0.]])
```

Derived Features

- Anything that is a transformation of our data
- This is where the money is!

Polynomial Features

```
In [52]: sns.pairplot(df_taxi[['passenger_count', 'trip_duration', 'tip_amount']]);
```



Polynomial Features Continued

In [53]: `from sklearn.preprocessing import PolynomialFeatures`

```
pf = PolynomialFeatures(degree=2,  
                        include_bias=False)  
X_new = pf.fit_transform(df_taxi[['passenger_count', 'trip_duration']])  
  
new_columns = ['passenger_count', 'trip_duration', 'passenger_count^2', 'passenger_count*trip_duration', 'trip_duration^2']  
pd.DataFrame(X_new[3:5], columns=new_columns)
```

Out[53]:

	passenger_count	trip_duration	passenger_count^2	passenger_count*trip_duration	trip_duration^2
0	3.0	298.0	9.0	894.0	88804.0
1	1.0	396.0	1.0	396.0	156816.0

Natural Language Processing (NLP)

- Many NLP Tasks
 - topic modeling
 - sentiment analysis
 - entity recognition
 - machine translation
 - natural language generation
 - question answering
 - relationship extraction
 - automatic summarization
 - ...
- All depend on extracting features from unstructured text data

NLP: The Corpus

- **corpus:** collection of documents
- Each item a document
 - tweet
 - review
 - resume
 - book
 - article
 - sentence
 - ...

NLP: Doc Representation

```
In [54]: doc_1 = "The cat in the hat."  
        doc_2 = "The quick brown cat jumped over the lazy cat."  
        corpus = [doc_1, doc_2]
```

- **tokens:** strings that make up a document ('the','cat',...)
- **terms:** unique set of strings in a documents
- **vocabulary:** set of unique terms that can be in any document
- **tokenization:** process of transforming document into tokens

NLP: Tokenization

- common tokenization method: split on whitespace

```
In [55]: doc_1.split()
```

```
Out[55]: ['The', 'cat', 'in', 'the', 'hat.']
```

```
In [56]: doc_2.split()
```

```
Out[56]: ['The', 'quick', 'brown', 'cat', 'jumped', 'over', 'the', 'lazy', 'cat.']
```

- Additional transformations depend on problem:
 - lowercase
 - remove stopwords
 - **stemming**: reduce token to stem (eg: "tokenization" -> "tokeniz")
 - **lemmatization**: common form (eg: "tokenization" -> "tokenize")
 - and tags?
 - remove special characters

NLP: Doc Representation

- Most common representation: Bag of Words (BOW)
 - split document into tokens
 - ignore order (lose context!)

```
In [57]: sorted(doc_1.lower().replace('.', '').split())
```

```
Out[57]: ['cat', 'hat', 'in', 'the', 'the']
```

```
In [58]: sorted(doc_2.lower().replace('.', '').split())
```

```
Out[58]: ['brown', 'cat', 'cat', 'jumped', 'lazy', 'over', 'quick', 'the', 'the']
```

NLP: n-grams

- create new terms as combinations of n tokens
- captures local context
- vocabulary increases quickly

Unigrams: n = 1

```
In [59]: sorted(doc_1.lower().replace('.', '').split())
```

```
Out[59]: ['cat', 'hat', 'in', 'the', 'the']
```

Bigrams: n = 2

```
In [60]: tmp = ("<start> "+doc_1+" <end>").lower().replace('.', '').split()
sorted([(tmp[i]+'_'+tmp[i+1]) for i in range(len(tmp)-1)])
```

```
Out[60]: ['<start>_the', 'cat_in', 'hat_<end>', 'in_the', 'the_cat', 'the_hat']
```

NLP: Term Frequency

- **Term Frequency (TF):** Number of times a particular term occurs in a document

```
In [61]: tokens = []
         for doc in corpus:
             tokens.extend(doc.lower().replace('.', '').split())
         vocab = sorted(list(set(tokens)))
         vocab
```

```
Out[61]: ['brown', 'cat', 'hat', 'in', 'jumped', 'lazy', 'over', 'quick', 'the']
```

```
In [62]: termfreq = np.zeros((len(corpus), len(vocab)))
         for doc_idx, doc in enumerate(corpus):
             for term_idx, term in enumerate(vocab):
                 for token in doc.lower().replace('.', '').split():
                     if term == token:
                         termfreq[doc_idx, term_idx] += 1
         df_termfreq = pd.DataFrame(termfreq, index=['doc1', 'doc2'], columns=vocab)
         df_termfreq
```

```
Out[62]:
```

	brown	cat	hat	in	jumped	lazy	over	quick	the
doc1	0.0	1.0	1.0	1.0	0.0	0.0	0.0	0.0	2.0
doc2	1.0	2.0	0.0	0.0	1.0	1.0	1.0	1.0	2.0

NLP: Term Frequency

- **Document Frequency (DF):** Number of documents a term occurs in

```
In [63]: docfreq = df_termfreq.sum()  
docfreq.sort_values(ascending=False)
```

```
Out[63]: the      4.0  
cat       3.0  
quick     1.0  
over      1.0  
lazy      1.0  
jumped    1.0  
in        1.0  
hat       1.0  
brown     1.0  
dtype: float64
```


Stopwords

- **stopwords:** terms that (generally) have high DF and aren't informative
- ex: 'a', 'about', 'above', ...
- often removed prior to analysis

```
In [64]: stopwords = ['the', 'a', 'an']

tokens = []
for doc in corpus:
    tokens.extend(doc.lower().replace('.', '').split())
vocab = sorted(list(set(tokens)))
vocab = [x for x in vocab if x not in stopwords]
vocab

Out[64]: ['brown', 'cat', 'hat', 'in', 'jumped', 'lazy', 'over', 'quick']
```

NLP: CountVectorizer

```
In [65]: from sklearn.feature_extraction.text import CountVectorizer

cv = CountVectorizer(stop_words=None,                # can use 'english', but arguments against: https://scikit-learn.org/stable/
                    ngram_range=(1,1),              # only unigrams
                    token_pattern= r'(?u)\b\w\w+\b', # at least one word-character surrounded by boundaries
                    min_df=1,                        # has to occur in at least one document
                    max_df=1.0,                      # can occur in at most 100% of the documents
                    )
X = cv.fit_transform(corpus)
cv.vocabulary_
```

```
Out[65]: {'the': 8,
          'cat': 1,
          'in': 3,
          'hat': 2,
          'quick': 7,
          'brown': 0,
          'jumped': 4,
          'over': 6,
          'lazy': 5}
```

```
In [66]: X
```

```
Out[66]: <2x9 sparse matrix of type '<class 'numpy.int64'>'
         with 11 stored elements in Compressed Sparse Row format>
```

```
In [67]: X.todense() # term frequencies
```

```
Out[67]: matrix([[0, 1, 1, 1, 0, 0, 0, 0, 2],
                 [1, 2, 0, 0, 1, 1, 1, 1, 2]])
```

NLP: Tf-Idf

- What if some terms are still uninformative?
- Can we downweight terms that are in many documents?
- **Term Frequency - Inverse Document Frequency (TfIdf)**
 - $\text{tfidf}(t, d) = \text{tf}(t, d) * \text{idf}(t)$
 - $\text{idf}(t) = \log [n / \text{docfreq}(t)] + 1$

```
In [68]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
tfidf = TfidfVectorizer()  
X = tfidf.fit_transform(corpus)  
X.todense()
```

```
Out[68]: matrix([[0.          , 0.33425073, 0.46977774, 0.46977774, 0.          ,  
                  0.          , 0.          , 0.          , 0.66850146],  
                 [0.33241213, 0.47302794, 0.          , 0.          , 0.33241213,  
                  0.33241213, 0.33241213, 0.33241213, 0.47302794]])
```

NLP: Example 20Newsgroups

```
In [69]: from sklearn.datasets import fetch_20newsgroups
```

```
ngs = fetch_20newsgroups()
```

```
# grab 100 docs
```

```
docs = ngs['data'][:100]
```

```
docs[0]
```

```
Out[69]: "From: lerxst@wam.umd.edu (where's my thing)\nSubject: WHAT car is this!?\nNntp-Posting-Host: rac3.wam.umd.edu\nOrganization: U\niversity of Maryland, College Park\nLines: 15\n\nI was wondering if anyone out there could enlighten me on this car I saw\nthe other day. It was a 2-door sports car, looked to be from the late 60s/\nearly 70s. It was called a Bricklin. The doors were r\nreally small. In addition,\nthe front bumper was separate from the rest of the body. This is \nall I know. If anyone can tellme\na model name, engine specs, years\nof production, where this car is made, history, or whatever info you\nhave on this funky loo\nking car, please e-mail.\n\nThanks,\n- IL\n\n---- brought to you by your neighborhood Lerxst ----\n\n\n\n\n"
```

```
In [70]: cv = TfidfVectorizer(ngram_range=(1,2), #unigrams + bigrams
```

```
                        stop_words=None,
```

```
                        min_df=2,
```

```
                        max_df=.8
```

```
                    ).fit(docs)
```

```
X = cv.transform(docs)
```

```
X.shape
```


```
Out[70]: (100, 3894)
```

```
In [71]: list(cv.stop_words_)[:5]
```

```
Out[71]: ['hayward to', 'accidents in', 'or thought', 'list long', 'in they']
```

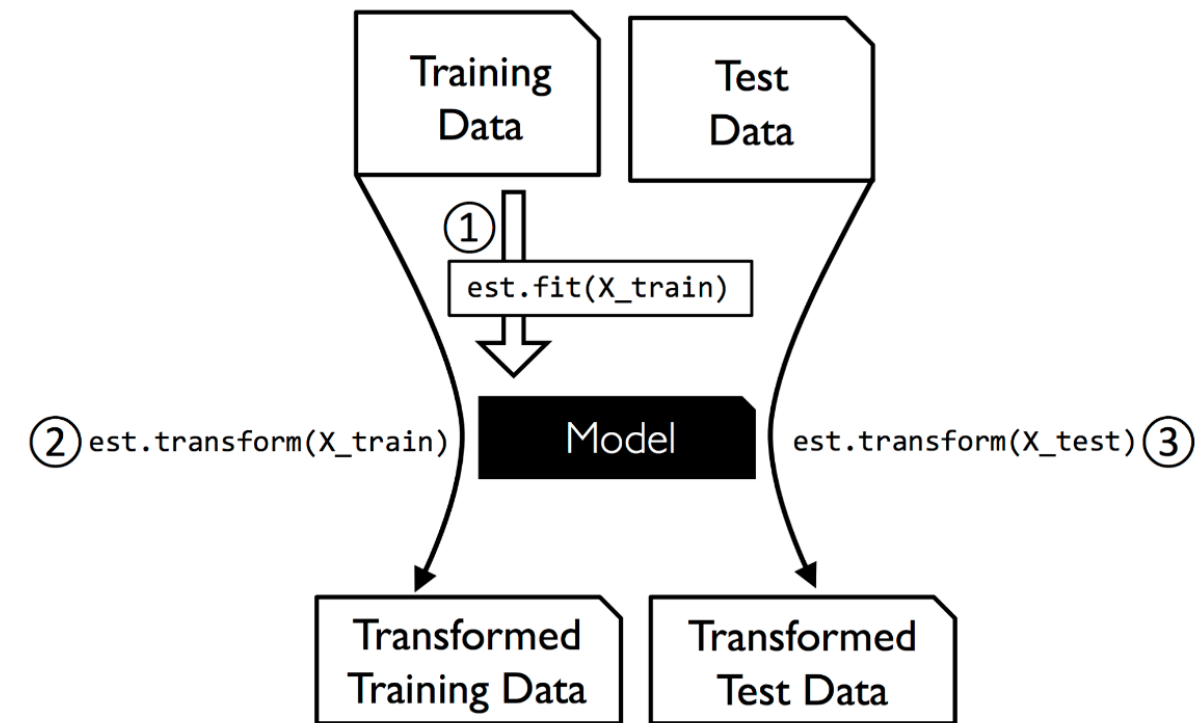
Predicting with Train/Test Split

- When training a model for prediction



Transforming with Train/Test Split

- When performing data transformation



Next

- Feature Selection
- Dimensionality Reduction
- Topic Modeling

Questions?